

Krishna Pillutla

Contact	Website: https://krishnap25.github.io Email: pillutla@cs.washington.edu	
Position	Visiting Researcher , Google Research	<i>Sept 2022 - Date</i>
Education	University of Washington Ph.D. in Computer Science & Engineering <i>Thesis:</i> From Enormous Structured Models to On-device Federated Learning: Robustness, Heterogeneity and Optimization <i>Advisors:</i> Zaid Harchaoui and Sham Kakade	<i>2016-2022</i>
	Carnegie Mellon University M.S. in Computer Science (QPA: 3.95/4.00) <i>Thesis:</i> Data Driven Resource Allocation for Distributed Learning <i>Advisor:</i> Maria-Florina Balcan	<i>2014-15</i>
	Indian Institute of Technology, Bombay B.Tech (Hons) in Computer Science & Engineering (QPA: 9.54/10.0) <i>Thesis:</i> Distributed Machine Learning: Iterative Convex Optimization Methods <i>Advisor:</i> J. Saketha Nath	<i>2010-14</i>
Awards	Outstanding Paper at NeurIPS Top 6 of 9000 submissions	<i>2021</i>
	J.P. Morgan PhD Fellowship 1 of 14 awardees worldwide	<i>2019-20</i>
	Anne Dinning - Michael Wolf Endowed Regental Fellowship First-year PhD Fellowship awarded on merit	<i>2016-17</i>
	CBSE Merit Scholarship by the Central Board of Secondary Education in India Awarded by the Govt. of India for the duration of undergraduate studies	<i>2010-14</i>
Previous Positions	Research Intern , Facebook AI Research <i>Host:</i> Lin Xiao <i>Project:</i> Federated learning with partial model personalization	<i>2021</i>
	Research Intern , Facebook AI Research <i>Hosts:</i> Awni Hannun, Gabriel Synnaeve <i>Project:</i> Sequence level end-to-end training of speech recognition models.	<i>2019</i>
	Machine Learning Engineer Intern , Rocket Fuel Inc. <i>Project:</i> Large-scale recommender systems to provide personalized advertisement recommendations, on both model-building and systems.	<i>2016</i>
	Research Intern , Microsoft Research <i>Hosts:</i> Dhruv Mahajan, Sundararajan Sellamanickam <i>Project:</i> Iterative parameter averaging for distributed machine learning	<i>2013</i>

Publications

Working papers and manuscripts¹:

- **Pillutla, K.**, Roulet, V., Kakade, S. M., & Harchaoui, Z. (2022)
The Trade-offs of Incremental Linearization Algorithms for Nonsmooth Composite Problems.
Submitted.
- Mehta, R., Roulet, V., **Pillutla, K.**, Liu, L. & Harchaoui, Z. (2022)
Stochastic Algorithms for Ordered Empirical Risk Minimization.
Submitted.
- Fisher, J., Liu, L., **Pillutla, K.**, Choi, Y., Harchaoui, Z. (2022)
Statistical and Computational Guarantees for Influence Diagnostics.
Submitted.
- Fisher, J., Jiang, L., **Pillutla, K.**, Swayamdipta, S., Hessel, J., Harchaoui, Z., Choi, Y. (2022)
On Model Revision: Content Re-mapping and Avoidance.

Peer-reviewed journal and conference papers:

- **Pillutla, K.***, Laguel, Y.*, Malick, J., & Harchaoui, Z. (2022)
Federated Learning with Superquantile Aggregation for Heterogeneous Data.
Machine Learning (to appear).
- **Pillutla, K.**, Malik, K., Mohamed, A., Rabbat, M., Sanjabi, M., & Xiao, L. (2022).
Federated Learning with Partial Model Personalization.
International Conference on Machine Learning (ICML).
- **Pillutla, K.**, Kakade, S. M., & Harchaoui, Z. (2022).
Robust Aggregation for Federated Learning.
IEEE Transactions on Signal Processing.
- **Pillutla, K.**, Swayamdipta, S., Zellers, R., Thackstun, J., Welleck, S., Choi, Y. & Harchaoui, Z. (2021).
MAUVE: Measuring the Gap Between Machine Text and Human Text using Divergence Frontiers.
Neural Information Processing Systems (NeurIPS). **Outstanding Paper Award (Top 6 of 9000)**.
- Liu, L., **Pillutla, K.**, Welleck, S., Oh, S., Choi, Y. & Harchaoui, Z. (2021).
Divergence Frontiers for Generative Models: Sample Complexity, Quantization Effects, and Frontier Integrals.
Neural Information Processing Systems (NeurIPS).
- Kusupati, A., Wallingford, M., Ramanujan, V., Somani, R., Park, J. S., **Pillutla, K.**, Jain, P., Kakade, S., & Farhadi, A. (2021).
LLC: Accurate, Multi-purpose Learnt Low-dimensional Binary Codes.
Neural Information Processing Systems (NeurIPS).
- Laguel, Y., **Pillutla, K.**, Malick, J., & Harchaoui, Z. (2021).
Superquantiles in Action: Subdifferential Calculus in Practice and Applications in Machine Learning.
Set Valued and Variational Analysis.
- Laguel, Y.*, **Pillutla, K.***, Malick, J., & Harchaoui, Z. (2021).
A Superquantile Approach to Federated Learning with Heterogeneous Devices.
IEEE Conference on Information Sciences and Systems (CISS).
- **Pillutla, K.**, Roulet, V., Kakade, S. M., Harchaoui, Z. (2018).
A Smoother Way to Train Structured Prediction Models.
Neural Information Processing Systems (NeurIPS).
- Jain, P., Kakade, S. M., Kidambi, R., Netrapalli, P., **Pillutla, V. K.**, & Sidford, A. (2017).
A Markov Chain Theory Approach to Characterizing the Minimax Optimality of Stochastic Gradient Descent (for Least Squares).
Foundations of Software Technology and Theoretical Computer Science (FSTTCS).

¹equal contribution denoted by * and alphabetical order by α

- Ruffalo, M., Stojanov, P., **Pillutla, V. K.**, Varma, R., & Bar-Joseph, Z. (2017). Reconstructing cancer drug response networks using multitask learning. *BMC Systems Biology*.
- Dick, T.^α, Li, M.^α, **Pillutla, V. K.**^α, White, C.^α, Balcan, M-F., & Smola, A. (2017). Data Driven Resource Allocation for Distributed Learning. *Artificial Intelligence and Statistics Conference (AISTATS)*.
- **Pillutla, V. K.**^{*}, Fang, Z.^{*}, Devineni, P., Faloutsos, C., Koutra, D., & Tang, J. (2016). On Skewed Multi-dimensional Distributions: the FusionRP Model, Algorithms, and Discoveries. *SIAM International Conference on Data Mining*.

Workshop papers:

- **Pillutla, K.**^{*}, Laguel, Y.^{*}, Malick, J., & Harchaoui, Z. (2022). Tackling Distribution Shifts in Federated Learning with Superquantile Aggregation. *NeurIPS 2022 Workshop on Distribution Shifts*. **Spotlight Presentation**.
- **Pillutla, K.**^{*}, Laguel, Y.^{*}, Malick, J., & Harchaoui, Z. (2022). Differentially Private Federated Quantiles with the Distributed Discrete Gaussian Mechanism. *International Workshop on Federated Learning: Recent Advances and New Challenges (FL-NeurIPS)*.
- **Pillutla, K.**, Roulet, V., Kakade, S. M., & Harchaoui, Z. (2022). The Trade-offs of Incremental Linearization Algorithms for Nonsmooth Composite Problems. *Order up! The Benefits of Higher-Order Optimization in Machine Learning (HOO-NeurIPS)*.
- **Pillutla, K.**, Kakade, S. M., & Harchaoui, Z. (2020). Robust Aggregation for Federated Learning. *International Workshop on Federated Learning for User Privacy and Data Confidentiality (FL-ICML)*. **Long Presentation**.
- Dick, T.^α, Li, M.^α, **Pillutla, V. K.**^α, White, C.^α, Balcan, N., & Smola, A. (2017). Data Driven Resource Allocation for Distributed Learning. *AAAI Workshop on Distributed Machine Learning*.

Software Released

Mauve: Measuring the Gap Between Neural Text and Human Text

- A package to compute the Mauve score for neural text generation
- Install as `pip install mauve-text`: **5000 monthly downloads**
- Software is available at <https://github.com/krishnap25/mauve> and documentation at <https://krishnap25.github.io/mauve/>
- Also available via the HuggingFace Evaluate package

Geom-Median: Fast and Differentiable Geometric Median in PyTorch and NumPy

- Implementation of the smoothed Weiszfeld algorithm to compute the geometric median, a high-dimensional median analogue
- Install as `pip install geom-median`: **265 monthly downloads**
- Software is available at https://github.com/krishnap25/geom_median

SQwash: Distributionally robust learning in PyTorch with a 1 additional line of code

- Implementation of the superquantile to add distributional robustness to PyTorch training loops
- Install as `pip install sqwash`: **65 monthly downloads**
- Software is available at <https://github.com/krishnap25/sqwash> and documentation at <https://krishnap25.github.io/sqwash/>

Casimir: Catalyst, smoothing, and inference.

A toolbox of selected optimization algorithms including Casimir-SVRG (as well as SVRG and Catalyst-SVRG as special cases) for unstructured tasks such as binary classification, and structured prediction tasks such as named entity recognition and visual object localization. Available at <https://github.com/krishnap25/casimir>.

RFA: Robust Federated Aggregation

Implementation of robust aggregation for federated learning in *TensorFlow Federated*. Available at this URL.

Invited Talks

Federated Learning with Partial Model Personalization (2022).
Federated Learning One World Seminar.

Federated Learning with Superquantile Aggregation for Heterogeneous Data (2021-22).
IFDS Ethics and Algorithms, International Conference on Continuous Optimization.

From Enormous Structured Models to On-device Federated Learning: Robustness, Heterogeneity, and Optimization (2022).

Microsoft Research, Meta AI Research, Google Research.

MAUVE: Measuring the Gap Between Neural Text and Human Text (2022).

Stanford NLP Seminar, Microsoft Research Asia, IFML NSF Site Visit.

Mentoring

- Mentoring two junior graduate students in research 2021-date
 - Jillian Fisher
 - Ronak Mehta
- Undergraduate Student Mentor: Mentored 12 freshmen in their overall development as they deal with academic pressures and transition into college life 2013-14
- Department Academic Mentor: Served as a liaison between students struggling with academics and professors 2012-14

Teaching

Statistical Learning with Differentiable Programming Winter 2021 and 2022

- Co-taught a class of 25-30 students in the professional Masters in Data Science program
- Co-designed the syllabus; prepared and delivered the “lab” component of each lecture; designed homework and exams; co-managed the logistics; held office hours and graded

Machine Learning for Big Data

Spring 2018

- Teaching assistant for 50 early-stage PhD students
- Prepared data and designed homework; co-designed, prepared, advised on course projects; designed exams; held office hours and graded

Reinforcement Learning and Bandits, Teaching Assistant

2019

Algorithms and Foundation of Computing, Volunteer Tutor

2016-17

Programming 101, Chemistry 101, Numerical Analysis, Teaching Assistant at IITB

2012-14

Academic Honors

- Perfect 100 percentile (top 8 out of 174,000) in Common Admission Test (CAT) 2013
- Gold medal at the Indian National Chemistry Olympiad (INChO). Part of initial shortlist for the International Chemistry Olympiad (Top 35 from 28,000) 2010
- Secured All India Rank 22 in IITJEE, an exam taken by half million students 2010
- Awarded the Certificates of Merit by the CBSE ² for being in the top 0.1% in India in Mathematics and Chemistry in Grade 12 examinations, AISSCE 2010

²CBSE is the Central Board of Secondary Education in India

Service

- **Workshop Co-organizer:** IFDS Workshop on Distributional Robustness in Data Science (2022)
- **Reviewer** for JMLR, Math. Prog., NeurIPS, AISTATS, JOTA, AISTATS, ICLR
- **Student Area Chair** for Machine Learning, UW CSE Graduate Admissions (2020-21)
- **Organizer** for New Graduate Student Orientation at UW (2017) and Panelist (2018-20)

**Relevant
Coursework**

Interactive Machine Learning, Convex Analysis and Nonsmooth Optimization, Deep Learning Systems, Reinforcement Learning and Bandits, Advanced Probability, Randomized Algorithms, Intermediate Statistics, Convex Optimization, Natural Language Processing, Computer Vision, Foundations of Machine Learning and Data Science, Probabilistic Graphical Models, Statistical Machine Learning, Advanced Introduction to Machine Learning, Data Mining